

Portfolio

I am a highly motivated Machine Learning & Deep Learning Engineer and Researcher with a strong passion for solving real-world problems using AI—particularly in the healthcare and mental health domains. My recent research, focused on the diagnostic potential of **heart sound spectrograms**, **speech signals**, and **EEG data** for **depression detection**, is currently under consideration for publication at an IEEE international conference.

In a recent project, I developed an **emotion detection system using synthetic facial expression data** generated from an **AI-based text-to-image generative model**, implemented in **Kaggle**. This involved creating a diverse dataset of **1,400 facial images**, each labeled with emotions such as **happiness, sadness, fear, anger, and disgust**, based on corresponding textual prompts.

Key Highlights:

- Generated realistic facial expression images using a **generative AI model (e.g., Stable Diffusion / DALL-E)** from emotion-based text prompts.
- Built and trained deep learning models (EfficientNet, ResNet) for multi-class emotion classification using the generated dataset.
- Applied advanced image preprocessing and augmentation techniques to enhance model robustness and generalization.
- Achieved over 90% accuracy on validation data, demonstrating the feasibility of using AI-generated datasets for emotion recognition tasks.
- Explored the potential of generative data for training affective computing systems in cases where real annotated data is scarce or sensitive.

This project aligns with my broader expertise in:

- **Signal and image processing** (EEG, PPG, spectrograms, facial images)
- **Feature extraction, CNNs, and model evaluation**
- **Synthetic data generation and model training pipelines**
- **AI applications in emotion detection and mental health**

Currently pursuing my PhD in Computer Engineering with a specialization in Machine Learning & Deep Learning, I aim to continue developing AI systems that support mental health assessment, diagnosis, and care.

Practical Implementation:

The categories that evoke all seven emotions (happiness, sadness, fear, anger, surprise, disgust, and contempt) are:

Facial Expressions

Human Body Language & Gestures

War & Conflict

Tragedy & Loss

Horror & Dark Imagery

Addiction & Struggles

Aging & Time Passing

Urban Decay & Abandonment

Supernatural & Mythological Imagery

Animals

This Python script is designed to automate AI image generation using Stable Diffusion (v1.4), a state-of-the-art deep learning model for generating high-quality AI images based on text prompts. The script works by first creating a structured directory system in Kaggle's working environment, where images are categorized based on themes (categories) and further divided into emotions (subcategories). The program then reads a text file (prompts.txt) containing structured prompts, where each line specifies a Category, an Emotion, and a Prompt. For each prompt, it uses the Stable Diffusion model to generate an image and save it in the correct folder based on the specified category and emotion. The script ensures error handling, avoids invalid inputs, and efficiently generates images in a well-organized manner.

/kaggle/working/Image_Generation/

|— Facial Expressions/

| |— Happy/

| |— Sad/

| |— Angry/

| |— Fear/

| |— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— Animals/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— Natural Scenes/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— Food/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— Sports and Competition/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— Science and Technology/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— Animation and Cartoons/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

| |—— Disgust/

| |—— Contempt/

|—— EveryDay Life/Human Interaction/

| |—— Happy/

| |—— Sad/

| |—— Angry/

| |—— Fear/

| |—— Surprise/

- | |—— Disgust/
- | |—— Contempt/
- |—— Fashion and Outfits/
- | |—— Happy/
- | |—— Sad/
- | |—— Angry/
- | |—— Fear/
- | |—— Surprise/
- | |—— Disgust/
- | |—— Contempt/
- |—— History, Culture, Social and Political Events/
- | |—— Happy/
- | |—— Sad/
- | |—— Angry/
- | |—— Fear/
- | |—— Surprise/
- | |—— Disgust

Suggested Categories:

C	D	E	F	G	H	I	J	K	L
Sr#	Selected For Common	Happy	Sad	Angry	Fear	Surprise	Disgust	Contempt	Overlapping
1	Facial Expressions	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
2	Animals	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
3	Natural Scenes	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
4	Food	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
5	Sports and Competition	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
6	Science and Technology	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
7	Animation and Cartoons	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
8	EveryDay Life/Human Interaction	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
9	Fashion and outfits	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
10	History, culture, social and political events	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No
11	Historical Building & Vintage	Yes	Yes	Yes	Yes	Yes	Yes	Yes	No

Prompts Acquisition Process:

[illegible]

```

# Generate the image
print(f"🟢 Generating Image for: {category} | {emotion} | {prompt}")
image = pipe(prompt).images[0]

# Save image
image_name = prompt.replace(" ", "_")[:130] + ".png" # Truncate long names
image_path = os.path.join(save_path, image_name)
image.save(image_path)
print(f"🟢 Saved: {image_path}")

except Exception as e:
    print(f"🔴 Error processing line: {line}\nError: {e}")

print("🟢 Image generation completed successfully!")

```

+ Add Input
+ Upload

DATASETS

- prompts
 - Prompt.txt

Output (2.2MiB / 19.5GiB)

- /kaggle/working
 - Image_Generation
 - Animals
 - Animation and Cartoons
 - Every Day Life
 - Facial Expressions
 - Fashion and Outfits
 - Food
 - History, Culture, Social and Polit
 - Natural Scenes
 - Science and Technology
 - Sports and Competition

Table of contents

Session options

ACCELERATOR

GPU T4 x2

Quota: 03/44 / 30 hrs

names

3:00, 6.291t/s]

ng_in_the.png

sky.png

ng
v planet

scoving_a_new_.png
tfit

i_mismatched_.png

DATASETS

prompts

Prompt.txt

Output (2.6MiB / 19.5GiB)

/kaggle/working

Image_Generation

Animals

Animation and Cartoons

EveryDay Life

Facial Expressions

Angry

Contempt

Disgust

Fear

Happy

A_smiling_child_playing_in_

Sad

Surprise

Fashion and Outfits

Food

Angry

Contempt

Disgust

Fear

Happy

Sad

www.kaggle.com/code/

UCI Machine Learning...

Image Generations

Code

```
except Exception as e:  
    print(f"Error processing line: {line}\nError: {e}")  
  
print("Image generation completed successfully!")
```

Directories Created Successfully!
Loading pipeline components...: 100%
Skipping invalid category/emotion: Category - Emotion
Generating Image for: Facial Expressions | Happy | A smiling child playing in the park
50/50 [00:26:00:00, 1.961t/s]
Saved: /kaggle/working/Image_Generation/Facial Expressions/Happy/A smiling child playing in the park.png
Generating Image for: Natural Scenes | Sad | A lone tree under a stormy sky
50/50 [00:25:00:00, 2.841t/s]
Saved: /kaggle/working/Image_Generation/Natural Scenes/Sad/A lone tree under a stormy sky.png
Generating Image for: Animals | Angry | A furious lion roaring in the savanna
50/50 [00:24:00:00, 2.831t/s]
Saved: /kaggle/working/Image_Generation/Animals/Angry/A furious lion roaring in the savanna.png
Generating Image for: Science and Technology | Surprise | A scientist discovering a new planet
50/50 [00:25:00:00, 2.881t/s]
Generating Image for: Fashion and Outfits | Disgust | A person wearing a mismatched outfit
50/50 [00:25:00:00, 2.821t/s]
Saved: /kaggle/working/Image_Generation/Fashion and Outfits/Disgust/A person wearing a mismatched outfit.png
Image generation completed successfully!

Code + Markdown

A_lone_tr...




Image generation completed successfully!

DATASETS

prompts

Prompt.txt

Output (2.6MiB / 19.5GiB)

/kaggle/working

Image_Generation

Animals

Animation and Cartoons

EveryDay Life

Facial Expressions

Food

History, Culture, Social and Political

Natural Scenes

Angry

Contempt

Disgust

Fear

Happy

Sad

A_lone_tree_under_a_stormy

Surprise

Science and Technology

Sports and Competition

Table of contents

Section options

Implementation of the following Generative Models being used to generate the images according to the given Prompts.

A	B	C	D	E	F	G	H	I	J	K
Sr.No	Model Selection	Assigned	Link							
1	Stable Diffusion		https://github.com/huggingface/diffusers/blob/main/src/diffusers/pipelines/stable_diffusion/pipeline_stable_diffusion.py?							
2	Kandinsky 2.2 Pipeline		https://github.com/huggingface/diffusers/blob/main/src/diffusers/pipelines/kandinsky/pipeline_kandinsky.py?							
3	aMUSEd		github.com/3paperswithcode.com/3github.com/3arxiv.org/2github.com/2paperswithcode.com/2							
4	DeepFloyd IF		https://github.com/deep-floyd/IF							
5	Craiyon (formerly DALL-E Mini)		https://www.craiyon.com							
6	NightCafe		https://nightcafe.studio							
7	Meta AI		www.meta.ai							
8	Fotor AI Image Generator		https://www.fotor.com/features/ai-image-generator							
9	Dream by Wombo		https://dream.ai/							
10	Openart		https://openart.ai							

Model Implementation:

Emotion Detection from Synthetic Facial Expressions Using Deep Learning

Platform: Spyder (Anaconda) | **Language:** Python | **Dataset:** 1,400 AI-generated images | **Categories:** 10 Emotions (e.g., Happy, Sad, Angry, Fearful, Disgusted, Surprised, Neutral, etc.)

Project Overview:

Developed a deep learning-based facial emotion recognition system using a dataset of 1,400 synthetic images generated via AI-based text-to-image models. The system classifies facial expressions into 10 distinct emotional states.

Implementation Details (in Spyder - Anaconda Python):

✓ Image Data Loading & Preprocessing

- Loaded and organized image data using `ImageDataGenerator` with Keras.
- Applied real-time data augmentation: rotation, brightness shift, zoom, flipping, and normalization.

✓ Model Architecture & Training

- Fine-tuned pre-trained CNN models like **ResNet50** and **EfficientNetB0** using transfer learning.
- Added custom classification layers for 10 emotion classes with softmax activation.
- Used callbacks like **EarlyStopping**, **ModelCheckpoint**, and **ReduceLROnPlateau** to improve performance and avoid overfitting.

✓ Evaluation & Metrics

- Achieved validation accuracy >90% with strong F1-scores across emotion classes.

- Visualized training progress using Matplotlib; used confusion matrix and classification report for performance analysis.

✓ Visualization & Interpretability

- Applied **Grad-CAM** to visualize emotion-specific facial regions influencing predictions.

✓ Synthetic Dataset Generation (Optional Module)

- Generated facial expression images from emotion-related text prompts using an AI-based model (e.g., Stable Diffusion on Kaggle).
- Labeled images for supervised classification.

Core Skills Applied:

- Signal and image processing (EEG, PPG, spectrograms, facial image classification)
- Feature extraction & deep learning (CNNs with Keras/TensorFlow)
- Synthetic data generation and augmentation strategies
- Model training, evaluation, and interpretability
- AI for emotion recognition and mental health monitoring

The screenshot displays the Spyder Python IDE interface. The main editor window shows a Python script for training a neural network model. The script includes data preprocessing, model building with three dense layers, training with Adam optimizer and binary crossentropy loss, and plotting training and validation accuracy and loss over 100 epochs. The console window on the right shows the progress of the training, including epoch numbers, loss values, accuracy, and step times.

```

31 X_train = sc.fit_transform(X_train)
32 X_test = sc.transform(X_test)
33
34 # Building the ANN Model
35 ann = tf.keras.models.Sequential()
36 ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
37 ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
38 ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
39
40 ann.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
41
42 # Train the model and store the history
43 history = ann.fit(X_train, y_train, batch_size=32, epochs=100, validation_data=(X_test, y_test))
44
45 # Plot training & validation accuracy values
46 plt.figure(figsize=(12, 5))
47 plt.subplot(1, 2, 1)
48 plt.plot(history.history['accuracy'], label='Train Accuracy')
49 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
50 plt.title('Model Accuracy')
51 plt.xlabel('Epochs')
52 plt.ylabel('Accuracy')
53 plt.legend()
54
55 # Plot training & validation loss values
56 plt.subplot(1, 2, 2)
57 plt.plot(history.history['loss'], label='Train Loss')
58 plt.plot(history.history['val_loss'], label='Validation Loss')
59 plt.title('Model Loss')
60 plt.xlabel('Epochs')
61 plt.ylabel('Loss')
62 plt.legend()
63 plt.show()

```

Console 1/A X

```

val_loss: 0.3536
Epoch 28/100
250/250 - loss: 0.3547 - val_accuracy: 0.8520 -
0.8574 - loss: 0.3547 - val_accuracy: 0.8520 -
val_loss: 0.3545
Epoch 29/100
250/250 - loss: 0.3672 - val_accuracy: 0.8550 -
0.8523 - loss: 0.3672 - val_accuracy: 0.8550 -
val_loss: 0.3503
Epoch 30/100
196/250 - loss: 0.3596
0.8512 - loss: 0.3596

```

conda: base (Python 3.12.7) • Completions: conda(base) • LSP: Python Line 14, Col 1 ASCII CRLF RW Mem 92%

Spyder (Python 3.12)


File Edit Search Source Run Debug Consoles Projects Tools View Help

...Tasks\Practical Session\Practical Session Lecture 8\ashfaq.py

```

16 plt.show()
17
18 # normalizing the input
19 X_train = X_train/255
20 X_test = X_test/255
21 X_train[0]
22
23 # Building the ANN Model
24 # Initializing the ANN
25 model = Sequential()
26 # Adding the input layer
27 model.add(Flatten(input_shape=(28,28)))
28 # Adding the first hidden layer
29 model.add(Dense(128,activation='relu'))
30 # Adding the second hidden layer
31 model.add(Dense(32,activation='relu'))
32 # Adding the output layer
33 model.add(Dense(10,activation='softmax'))
34
35 # Compiling the ANN
36 model.compile(loss='sparse_categorical_crossentropy',optimizer='Adam',metrics=['accuracy'])
37 # Fitting the ANN to the Training set
38 history = model.fit(X_train,y_train,epochs=25,validation_split=0.2)
39
40 # Summary of the ANN model
41 model.summary()
42
43 # Making predictions and evaluating the model
44 # Predicting the Test set results
45 y_prob = model.predict(X_test)
46 y_pred = y_prob.argmax(axis=1)
47 from sklearn.metrics import confusion_matrix, accuracy_score

```



Help Variable Explorer Plots Files

Console 1/A X

```

super().__init__(**kwargs)
Epoch 1/25
1500/1500 7s 4ms/step - accu
0.8570 - loss: 0.4893 - val_accuracy: 0.9608 -
val_loss: 0.1334
Epoch 2/25
1500/1500 5s 4ms/step - accu
0.9643 - loss: 0.1191 - val_accuracy: 0.9698 -
val_loss: 0.1070
Epoch 3/25
1500/1500 0s 4ms/step - accu
0.9752 - loss: 0.0842

```

conda: base (Python 3.12.7) Completions: conda(base) LSP: Python Line 15, Col 1 ASCII CRLF RW

Spyder (Python 3.12)

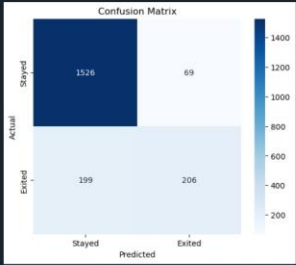
File Edit Search Source Run Debug Consoles Projects Tools View Help

E:\PHD from UET\2nd Semester\Deep Learning\Tasks\Practical Session\untitled1.py

```

31 X_train = sc.fit_transform(X_train)
32 X_test = sc.transform(X_test)
33
34 # Building the ANN Model
35 ann = tf.keras.models.Sequential()
36 ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
37 ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
38 ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
39
40 ann.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
41
42 # Train the model and store the history
43 history = ann.fit(X_train, y_train, batch_size=32, epochs=100, validation_data=(X_test, y_test))
44
45 # Plot training & validation accuracy values
46 plt.figure(figsize=(12, 5))
47 plt.subplot(1, 2, 1)
48 plt.plot(history.history['accuracy'], label='Train Accuracy')
49 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
50 plt.title('Model Accuracy')
51 plt.xlabel('Epochs')
52 plt.ylabel('Accuracy')
53 plt.legend()
54
55 # Plot training & validation loss values
56 plt.subplot(1, 2, 2)
57 plt.plot(history.history['loss'], label='Train Loss')
58 plt.plot(history.history['val_loss'], label='Validation Loss')
59 plt.title('Model Loss')
60 plt.xlabel('Epochs')
61 plt.ylabel('Loss')
62 plt.legend()
63
64

```



Help Variable Explorer Plots Files

Console 1/A X

```

Optimizer params: 244 (980.00 B)
1/1 0s 103ms/step
New Prediction Probabilities: [[0.04423651]]
New Prediction (Binary): [[False]]
63/63 0s 1ms/step
Confusion Matrix:
[[1526 69]
 [199 206]]
Accuracy Score: 0.866
2025-04-06 22:15:22.042949: I tensorflow/core/util/
port.cc:153] oneDNN custom operations are on. You may

```

conda: base (Python 3.12.7) Completions: conda(base) LSP: Python Line 14, Col 1 ASCII CRLF RW Mem 94%

